


# ECE 1175 Embedded Systems Design

## Lab 5 – CPU Frequency Control



# ECE 1175 – Lab 5


- **Manipulating CPU Frequency on Raspberry Pi**
  - Access CPU frequency info
  - Manually change CPU frequency
  - Task 1: warm-up of CPU frequency manipulation
- **Customize CPU Frequency Governor**
  - “schedutil” governor
  - Task 2: implement “schedutil” from userspace

# Prerequisites

## ■ CPU Frequency vs. Power Consumption

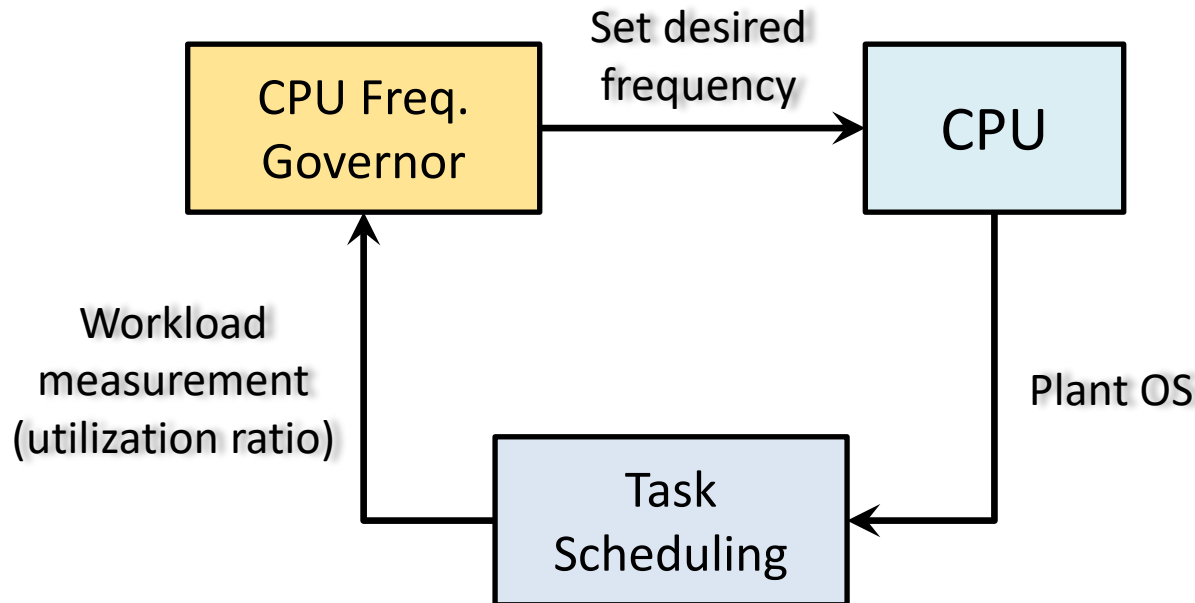
Dynamic CPU power consumption is proportional to the CPU frequency

$$P_{dyn} = CV^2f$$



# Prerequisites

- **CPU Frequency Governor**
  - Control CPU frequency based on different workloads
  - Basic idea: heavy load → high freq. Light load → low freq.



# Access CPU Frequency Info

## ■ Sysfs

- Linux kernel provides an interface via sysfs pseudo filesystem
- We can get access under:  
/sys/devices/system/cpu/cpu[#No.]/cpufreq/

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0 $ cd cpufreq
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ ls
affected_cpus          related_cpus           scaling_governor
cpuinfo_cur_freq       scaling_available_frequencies scaling_max_freq
cpuinfo_max_freq       scaling_available_governors scaling_min_freq
cpuinfo_min_freq       scaling_cur_freq       scaling_setspeed
cpuinfo_transition_latency scaling_driver          stats
```

# Access CPU Frequency Info

## ■ Useful CPU Frequency Info

- Check the maximum, minimum, and current CPU frequency
  - `cpuinfo_max_freq`
  - `cpuinfo_min_freq`
  - `cpuinfo_cur_freq`

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat cpuinfo_max_freq
1500000
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat cpuinfo_min_freq
600000
```

Note: the frequency unit is **kHz** by default

# Access CPU Frequency Info

## ■ Useful CPU Frequency Info

- Check available governors and frequencies
  - `scaling_available_governors`
  - `scaling_available_frequencies`

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat scaling_available_governors
conservative ondemand userspace powersave performance schedutil
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat scaling_available_frequencies
600000 750000 1000000 1500000
```

# Change CPU Frequency

- We can manually change CPU frequency to any **available** values. To do so,
  - Switch the governor to the **userspace** mode
  - Write your frequency value to the **scaling\_setspeed** file

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ sudo su 1. Switch to root
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo userspace > scaling_governor 2. Switch to userspace
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# cat scaling_governor
userspace
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# cat cpuinfo_cur_freq
600000
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo 1000000 > scaling_setspeed 3. Write any available frequency values
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# cat cpuinfo_cur_freq
1000000
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0#
```

Note: you must run these commands as **root**



# Lab Task 1

## ■ A Warm-up of CPU Frequency Manipulation

- Measure runtime of **matrix multiplication (lab3)** under different CPU frequencies. Pick an  $N$  of mm and test the program on two frequencies: *1.5GHz* and *600MHz*

```
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo userspace > scaling_governor
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo 1500000 > scaling_setspeed
root@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0# echo 600000 > scaling_setspeed
```

```
pi@raspberrypi:~ $ sudo taskset -c 0 ./lab5_1
Execution time: 2.670180s
pi@raspberrypi:~ $ sudo taskset -c 0 ./lab5_1
Execution time: 6.813416s
```

Note: Run your program on single core

## ■ Checkoff


- Show the runtime results of your program under the two CPU frequencies

# Schedutil Governor

## ■ Schedutil

- Already implemented in the kernel
- For details: <https://lkml.org/lkml/2016/3/17/420>

```
pi@raspberrypi:/sys/devices/system/cpu/cpu0/cpufreq $ cat scaling_available_governors  
conservative ondemand userspace powersave performance schedutil
```



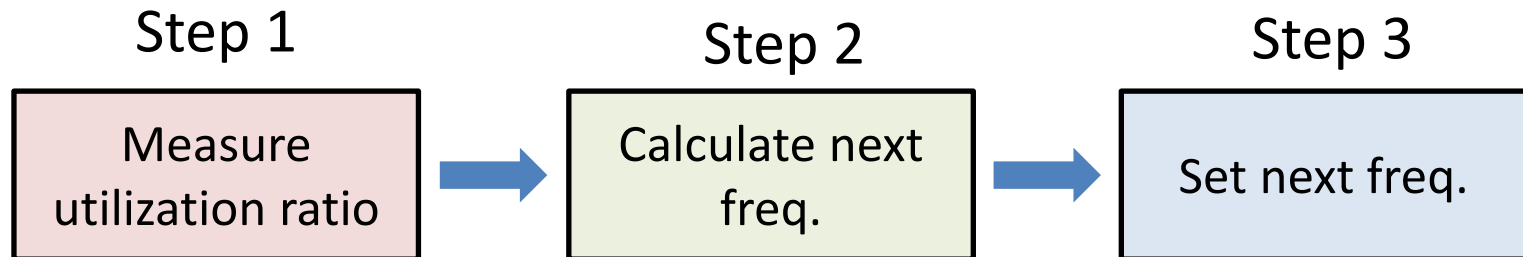
If utilization is frequency-invariant, schedutil says

$$\text{desired\_freq} = 1.25 * \text{max\_freq.} * \text{util\_ratio}$$

Next freq. is the available frequency **closest** to desired\_freq

# Lab Task 2

- Write **your own** program to implement schedutil from userspace



- We are very clear about how to do step 2 and 3
  - Step 2: just apply a formula
  - Step 3: write the value to scaling\_setspeed file

**But how to get utilization ratio in real-time?**

# Lab Task 2

- **To get real-time CPU utilization ratio**
  - Read from file: `/proc/stat`

```
pi@raspberrypi: /proc $ cat stat
```

cpu	51869	0	6607	4803123	4126	0	228	0	0	0
cpu0	28376	0	2787	1157054	1996	0	193	0	0	0
cpu1	8942	0	1250	1214468	635	0	22	0	0	0
cpu2	6293	0	1412	1216430	717	0	8	0	0	0
cpu3	8258	0	1158	1215171	776	0	5	0	0	0

```
intr 28774324 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 304386 0 0 0 0 0 6793 14 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 931 0 0 0 0 16619 0 0 16244 0 0 0 0 0 0 0 0 276  
56723 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 4545 5124 0 0 0 0 289289 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
ctxt 1743013  
btime 1604165995  
processes 3923  
procs_running 1  
procs_blocked 0  
softirq 842530 52919 213442 9 18098 0 0 270344 181986 0 105732
```

# Lab Task 2

## ■ Get per-CPU utilization ratio in real-time

- Interpret entries in `/proc/stat`
- For details: [Link](#)

For each `cpu[#No.]`,

**Idle time** since boot:

Sum of the **4** and **5** columns

```
cpu 51869 0 6607 4803123 4126 0 228 0 0 0
cpu0 28376 0 2787 1157054 1996 0 193 0 0 0
cpu1 8942 0 1250 1214468 635 0 22 0 0 0
cpu2 6293 0 1412 1216430 717 0 8 0 0 0
cpu3 8258 0 1158 1215171 776 0 5 0 0 0
```

**Total time** since boot:

Sum from **1** to **8** columns

The very first "cpu" line aggregates the numbers in all of the other "cpuN" lines. These numbers identify the amount of time the CPU has spent performing different kinds of work. Time units are in USER\_HZ (typically hundredths of a second). The meanings of the columns are as follows, from left to right:

- user: normal processes executing in user mode
- nice: niced processes executing in user mode
- system: processes executing in kernel mode
- idle: twiddling thumbs
- iowait: In a word, iowait stands for waiting for I/O to complete. But there are several problems:
  1. Cpu will not wait for I/O to complete, iowait is the time that a task is waiting for I/O to complete. When cpu goes into idle state for outstanding task io, another task will be scheduled on this CPU.
  2. In a multi-core CPU, the task waiting for I/O to complete is not running on any CPU, so the iowait of each CPU is difficult to calculate.
  3. The value of iowait field in `/proc/stat` will decrease in certain conditions.So, the iowait is not reliable by reading from `/proc/stat`.
- irq: servicing interrupts
- softirq: servicing softirqs
- steal: involuntary wait
- guest: running a normal guest
- guest\_nice: running a niced guest

# Lab Task 2

- **Get per-CPU utilization ratio in real-time**
  - Calculate utilization ratio for each CPU core

```
// Get accumulative t_total & t_usage
```

$$t_{total} = user + nice + system + idle + iowait + irq + softirq + steal$$
$$t_{idle} = idle + iowait$$
$$t_{usage} = t_{total} - t_{idle}$$

```
// Compute util ratio within a short period (from moment  $\tau_1$  to  $\tau_2$ )
```

```
// If  $\tau_1$  is the current timestep,  $\tau_2$  should be  $\tau_1 + \Delta t$  (e.g.,  $\Delta t = 0.5$  sec.)
```

$$\Delta t_{total} = t_{total}(\tau_2) - t_{total}(\tau_1)$$
$$\Delta t_{usage} = t_{usage}(\tau_2) - t_{usage}(\tau_1)$$
$$\%Util = \frac{\Delta t_{usage}}{\Delta t_{total}} \times 100\%$$

One more step: find the **highest** util ratio across all cores as the final util ratio fed to the formula

# Lab Task 2

- **Use programming language of your choice**
  - C/C++/Python/Others
- **Checkoff**
  - Demonstrate results following these steps
    - Switch to userspace mode
    - Keep running your governor program in the background
    - Run test programs with different workloads
    - Check how CPU frequency changes
  - Submit your code on Canvas

# Lab Task 2

1. First run your governor  
(use sudo or run in root)

2. Then run your test program  
(e.g., your task1 program)

3. Check current frequency  
(Or you can integrate this part  
into your governor program)



# Lab Task 2

```
pi@raspberrypi: ~  
File Edit Tabs Help  
CPU0: 11.211111%  
CPU1: 0.100000%  
CPU2: 22.322222%  
CPU3: 0.100000%  
max util: 22.322222%  
next freq: 600000
```

```
pi@raspberrypi: ~  
File Edit Tabs Help  
pi@raspberrypi:~ $ ./lab5_1  
Execution time: 2.766234s pi@raspberrypi:~ $
```

```
pi@raspberrypi: /sys/devices/system/cpu/cpufreq/policy0  
File Edit Tabs Help  
pi@raspberrypi:~ $ cd /sys/devices/system/cpu/cpufreq/policy0  
pi@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0 $ ls  
affected_cpus      related_cpus      scaling_governor  
cpuinfo_cur_freq   scaling_available_frequencies scaling_max_freq  
cpuinfo_max_freq   scaling_available_governors scaling_min_freq  
cpuinfo_min_freq   scaling_cur_freq  scaling_setspeed  
cpuinfo_transition_latency scaling_driver      stats  
pi@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0 $ cat scaling_cur_freq  
1500000  
pi@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0 $ cat scaling_cur_freq  
600000  
pi@raspberrypi:/sys/devices/system/cpu/cpufreq/policy0 $
```

busy

idle

You should observe different frequencies for running and not running the program

Thank you!

Have fun with your Raspberry Pi!